

## **Summary of Design & Documentation – Reinforcement Learning Puddle World Implementation**

### **Introduction to Application, Communication protocol & Architecture**

The application that I built is a client-server architecture-based reinforcement learning puddle world implementation. The puddle world is a reinforcement learning benchmark problem that involves a grid world where an agent is free to move around in, and it contains puddles or states that can significantly reduce the cumulative discounted reward achieved by the agent at the end of the reinforcement learning training process' episodes. The objective of the training process is teaching the reinforcement learning (RL) agent how to effectively navigate around the puddles better as it goes through more training episodes and maximizes the final discounted reward that it ends with (ending with a high positive cumulative episode discounted reward) in the final episode of training. The client-server architecture was adopted in this application to maintain a consistent architecture pattern as that of the Game of Amazons in the COSC 322 course here at UBC, as well as to delegate the action choosing and training logic to the client side for students of the COSC 322 course to complete, and the server to closely monitor the training process of the RL agent through a messaging protocol that enables full transparency during the RL agent's training process. The SmartFoxServer's Java server and client APIs were both used by my application's parts to communicate with each other.

The communication protocol for the client-server communication was defined as containing the following reinforcement-learning specific variables that were important to ensure that the training process is conducted as expected and in accordance with the Q-learning formula [1, see references]:

Client to server messages:

1. The current state of the RL agent (GAME\_STATE)
2. The available actions for an RL agent from the current state (GAME\_AVAILABLE\_ACTIONS)
3. The available rewards for the states that can be traversed from the current state (GAME\_AVAILABLE\_REWARDS)
4. The action chosen by the client side ((GAME\_ACTION\_MOVE)
5. The reward of the action taken by the RL agent ((GAME\_ACTION\_REWARD)
6. The client reaching the terminal state ((GAME\_FINAL\_STATE)
7. The client asking the server to reset the RL environment (GAME\_RESET)
8. The client sending Q table updates (GAME\_Q\_UPDATE)
9. The client sending V table updates (GAME\_V\_UPDATE)
10. The client asks for an episode summary (with the cumulative discounted reward and the termination condition) (GAME\_INFO)
11. The client sending a message to the server that the client side agent has finished training (GAME\_TRAINING\_COMPLETE)

Server to client responses:

1. The server responding with the reward of the action taken by the RL agent (GAME\_ACTION\_REWARD)
2. The server responding with the current state of the RL agent (GAME\_STATE\_RESPONSE)
3. The server responding with the available actions for the client side RL agent from the current state (GAME\_AVAILABLE\_ACTIONS\_RESPONSE)

4. The server responding with the available rewards for the states that can be traversed from the current state (GAME\_AVAILABLE\_REWARDS\_RESPONSE)
5. The server responding with the reward of the action taken by the client (GAME\_ACTION\_REWARD\_RESPONSE)
6. The server responding with notifying that the terminal state has been reached (GAME\_FINAL\_STATE\_RESPONSE)
7. The server responding with resetting the environment for the RL agent's training (GAME\_RESET\_RESPONSE)
8. The server responding with an episode summary (with the cumulative discounted reward and the termination condition) (GAME\_INFO\_RESPONSE)
9. The server responding with an error message based on incorrect parameters or format of the request sent by the client (GAME\_ERROR)

### **Setup & How to Run the Application**

The zipped code folder is called “deliverables 2”. Here are the dependencies that a user must have installed on their computer before unzipping the code:

1. SmartFoxServer 2X Java Server API in the Community tab:  
<https://www.smartfoxserver.com/download/sfs2x#p=installer> [2, see references]
2. SmartFoxServer 2X Java Client API in the Client API tab:  
<https://www.smartfoxserver.com/download/sfs2x#p=client> [3, see references]
3. Java: <https://www.oracle.com/ca-en/java/technologies/downloads/> [3, see references]
4. Maven
5. Dependencies in the “deliverable 2/ygraph-ai-smartfox-server/pom.xml” and “deliverable 2/ygraph-ai-smartfox-client/pom.xml”

Below is the process that must be followed to run the application:

1. Download and unzip the “deliverables 2” folder from any folder on your device.
2. Navigate to the directory titled “ygraph-ai-smartfox-client/target” and you will see the client side code from where the RL agent will be run from.
3. Open up your file system and look for this directory  
“/Applications/SmartFoxServer\_2X/SFS2X/extensions” or something similar, and copy the file titled “ygraph-ai-smartfox-server-extension-2.1.jar” in the “JAR file” directory in the deliverables 2 folder that you have open.
4. Create a directory called cosc322-2 in the  
“/Applications/SmartFoxServer\_2X/SFS2X/extensions” directory and insert the copied JAR file into it.
5. Now, go into the directory named “Zone file” in the “ygraph-ai-smartfox-client/target” directory and copy the file titled “cosc322-2.zone.xml”.
6. Next, go into the directory titled “/Applications/SmartFoxServer\_2X/SFS2X/zones” or something similar, and paste the copied file into this directory.
7. Navigate to the “deliverable 2/ygraph-ai-smartfox-server/pom.xml” file, and ensure that the text between <sfs.extensions.dir> </sfs.extensions.dir> has the path to your  
“/Applications/SmartFoxServer\_2X/SFS2X/extensions/” in it, and the text between <sfs.working.dir> </sfs.working.dir> has the path to your

“/Applications/SmartFoxServer\_2X/SFS2X/” or something similar to make sure the RL agent will run without any errors. Below is a section of my pom file:

```
<properties>
<maven.compiler.source>1.8</maven.compiler.source>
<maven.compiler.target>1.8</maven.compiler.target>

<sfs.extensions.dir>/Applications/SmartFoxServer_2X/SFS2X/extensions/</sfs.extensions.dir>
<sfs.working.dir>/Applications/SmartFoxServer_2X/SFS2X/</sfs.working.dir>
</properties>
```

8. Launch SFS2XLauncher on your device.
9. Now, open up a terminal and make sure that your current working directory “ygraph-ai-smartfox-client/target”, and run the following the command to run the RL agent’s training process:

```
java -jar ygraph-ai-smartfox-client-2.1.jar agent1 password1 localhost 9933 cosc322-2 RLRoom
```

10. Once you hit enter, you will see logs on the terminal you have open (client-side logs) and on the SFS2XLauncher app (server-side logs), which you can scroll up and down to monitor the training process of the RL agent.
11. Please take a look at the parameters on both client-side and server-side .env files in the “ygraph-ai-smartfox-client” and “ygraph-ai-smartfox-server” directories, which can be changed to change the RL agent’s training process’ termination condition, and if you choose to change the parameters in one .env file, please also change it in the other .env file.
12. Before running the code again after changing the parameters in both .env files, please comment out the entire test files called “/deliverable  
2/ygraph-ai-smartfox-server/src/test/java/ygraph/ai/smartfox/rl/test/java/ygraph/ai/smartfox/rl/test/RLGameMessageTest.java” and “/deliverable  
2/ygraph-ai-smartfox-server/src/test/java/ygraph/ai/smartfox/rl/test/java/ygraph/ai/smartfox/rl/test/RLWorldTest.java”.
13. Then, navigate to the “ygraph-ai-smartfox-server” directory and run the following command in your terminal:

```
mvn clean package
```

14. Once you have run the command, copy the file generated in the “deliverable  
2/ygraph-ai-smartfox-server/target” directory titled “ygraph-ai-smartfox-server-extension-2.1.jar” and paste it into the “/Applications/SmartFoxServer\_2X/SFS2X/extensions/cosc322-2” directory and replace the old JAR file that was there.
15. Now you can rerun the application and the modified parameters will be used to run the application.
16. It is necessary to repeat this process every time the parameters in the .env files are changed to ensure that the server can read the new values in the server side .env file.

### How does the application work?

The application, at a high-level, has a server-side that works to receive the chosen actions from the client-side, and use that to update its representation of the RL agent on the server-side, and check for validity of actions, move the RL agent, and keep the RL agent's training process going till a termination condition is hit. The server-side will respond to the client side with possible actions from a particular state, and their respective rewards, and the client can then use to decide on what action to choose based on the epsilon-greedy policy, that balances the ability of an RL agent to randomly choose an action (exploration) versus choosing the best action based on the computed Q-value in the Q-table (exploitation) [5, see references]. Q-learning is the process of training an RL agent to maximize its cumulative discounted reward based on Q-tables (states are rows, and actions are columns) that store estimates of how good it is to be in state  $s$  and take action  $a$  to maximize cumulative discounted rewards and are what guide the action suggestions of the RL agent from the client side. 4 possible actions are allowed for the RL agent: up, down, left, and right. I also defined the V-table, for which each box in the table tells us how much total discounted reward I will see being in state  $s$  in total. The following Q-value update formula was used [1, see references]:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]. \quad (6.8)$$

This update formula for Q-values in the Q-table guides the client side RL agent in making suggestions for actions, with  $S_t$  being the state that the RL agent is in at time interval  $t$ ,  $A_t$  being the action that the RL agent took at time interval  $t$ ,  $\alpha$  is the speed at which the RL agent learns to avoid puddles (it priorities immediate rewards over future rewards),  $Q(S_t, A_t)$  represents the Q-value of going with action  $A_t$  from state  $S_t$ ,  $R_{t+1}$  is the reward in time interval  $t + 1$ ,  $\gamma$  is the discount factor (it priorities future rewards over immediate rewards).

Each individual class on the client-side and server-side serve a different purpose in training the RL agent, at a high-level, they are explained below (please refer to the code files for further explanation of how the internal details of the code works):

Client-side classes:

1. RLClientGameMessage.java : This class creates and manages attributes of a message sent between the client and server - establishing a consistent protocol.
2. RLGameModel.java: This class holds the game's state specific to RL puddle world parameters like: state, available actions, rewards and the final state check.
3. RLGamePlayer.java: This class represents an RL agent's core logic and interactions with the Puddle World environment along with handling the connection and message processing to and from the server.

Server-side classes:

1. RLGameExtension.java: This class creates the RL environment within the server and sets up a request handler to handle client-related RL logic.
2. RLGameManager.java: This class takes RLGameUser instances and binds them with RLWorld instances allowing for concurrent usage of the application.
3. RLGameMessage.java: This class defines the messaging protocol that understands what is being sent from the client side and can communicate back to the client side.
4. RLGameRequestHandler.java: This class handles RL-specific client-side requests and sends responses back to the client.
5. RLGameUser.java: This class represents a user within an RL game (server-side representation of the RL agent).
6. RLMultiHandler.java: This class manages non-RL-specific client requests like joining rooms and using RLGameManager to add and remove users.
7. RLWorld.java: This class defines the internal representation of the Puddle world on the server side.

The control flow of the program from start to end is as follows:

The RLGameExtension registers the RLGameRequestHandler to handle the client-side's requests, and then RLGamePlayer on the client side logs in and joins the room via RLMultiHandler, and then RLGamePlayer starts by requesting the initial state, and the user added to the hashmap storing the RLGameUser objects mapped to their usernames. Next, the server responds with the state that the RL agent is in, and the client-side requests possible actions and their respective rewards, and the server responds to those. The client side picks the next action with the epsilon-greedy policy, updates the Q and V tables, and sends the Q and V updates to the server along with the action that it picked. This process keeps going till the termination conditions are reached, till the number of episodes are set as the parameters in the .env files for the client and server side. The 3 termination conditions for an episode are as follows based on the STOP\_METHOD variable in the .env files on the client and server side:

STOP\_METHOD=0:

- With this training termination condition, the RL agent is left to traverse the puddle world until the number of maximum steps defined in the client and server side .env files.
- It is advised to keep the MAX\_PUDDLES and PUDDLE\_SIZE small (if setting MAX\_PUDDLES=3 and PUDDLE\_SIZE=3 on gridSize=6, the number of states that get covered with puddles is  $3^3$  states (out of  $6^2$  total states on the grid), as the puddles are created in grouped squares, so there will be three 3x3 squares of puddle states, therefore, please either maintain MAX\_PUDDLES closer to 1 in smaller grid sizes, or if you prefer to keep the MAX\_PUDDLES and PUDDLE\_SIZE equal to each other, then maintain  $MAX\_PUDDLES^3 < \frac{1}{4}$ th of grid size in puddle states) as it can take a while to generate puddles that do not overlap each other or the goal state.
- This is the default behavior, and the training will keep going on until the maximum number of steps is reached or the goal state is reached, whichever comes first.

STOP\_METHOD=1:

- With this training termination condition, the RL agent is left to traverse the puddle world until it reaches the goal state, which is located at the state ID:  $gridSize^2 - 1$  (it is advised to keep the  $gridSize < 8$  or  $9$  to reduce the time it takes for a run of the program (n episodes of training to complete)).

STOP\_METHOD=2:

- With this training termination condition, the RL agent is left to traverse the puddle world until it is stopped by a probability that is randomly generated. This probability is set on the client and server side .env files, and essentially the way this method works is that a random number is generated, if it is lesser than the probability that is in the .env files (STOP\_PROB), then the episode is terminated immediately, but if the random number generated is greater than STOP\_PROB, then the training process moves to the next step, and a new random number is generated before each step is taken by the RL agent.
- This method will keep going on until either the probabilistic stopping terminates the episode, or if the RL agent reaches the goal state, whichever comes first.

### Output (Example Runs)

Here is what the .env files had in them when running the below tests of the application:

```
MAX_STEPS=30
EPISODE_COUNT=3
STOP_PROB=0.1
GRID_SIZE=5
ALPHA=0.1
GAMMA=0.9
EPSILON=1.0
DEFAULT_REWARD=-0.02
PUDDLE_REWARD=-2.0
GOAL_REWARD=20.0
MAX_PUDDLES=2
PUDDLE_SIZE=2
SUCCESS_REWARD_THRESHOLD=2.0
```

Snippets of the outputs on the client and server sides for each different stopping method are shown below, of which each visual has an episode summary to it, the number of steps in the episode, and the cumulative discounted reward from each episode (this output is focusing on a single episode of training due to the length of the logs):

STOP\_METHOD=0:

Client-side logs:

Episode 1:

```
Received EXTENSION_RESPONSE: cmd=rl.action
Received messageType: GAME_ACTION_REWARD_RESPONSE
Parsing messageType: GAME_ACTION_REWARD_RESPONSE
All received keys: [reward, messageType, nextStateId, action]
State transition: 17 -> 18
```

```
DEBUG: stopMethod=0
```

```
==== Episode 1 Information ====
```

```
Reason: Maximum steps (30) reached
```

```
Goal Reached: No
```

```
=====
```

```
Sent GAME_FINAL_STATE message to the server.
```

```
==== End of Episode Summary ====
```

```
Steps Taken: 30/30
```

```
Discounted Episode Reward: -8.811421888987708
```

```
=====
```

Episode 2:

```
Received EXTENSION_RESPONSE: cmd=rl.action
```

```
Received messageType: GAME_ACTION_REWARD_RESPONSE
```

```
Parsing messageType: GAME_ACTION_REWARD_RESPONSE
```

```
All received keys: [reward, messageType, nextStateId, action]
```

```
State transition: 17 -> 22
```

```
DEBUG: stopMethod=0
```

```
==== Episode 2 Information ====
```

```
Reason: Maximum steps (30) reached
```

```
Goal Reached: No
```

```
=====
```

```
Sent GAME_FINAL_STATE message to the server.
```

```
==== End of Episode Summary ====
```

```
Steps Taken: 30/30
```

```
Discounted Episode Reward: -11.213227822475048
```

```
=====
```

Episode 3:

```
Received EXTENSION_RESPONSE: cmd=rl.action
```

```
Received messageType: GAME_ACTION_REWARD_RESPONSE
```

```
Parsing messageType: GAME_ACTION_REWARD_RESPONSE
```

```
All received keys: [reward, messageType, nextStateId, action]
```

```
State transition: 15 -> 16
```

```
DEBUG: stopMethod=0
```

```
==== Episode 3 Information ====
```

```
Reason: Maximum steps (30) reached
```

```
Goal Reached: No
```

```
=====
```

```
Sent GAME_FINAL_STATE message to the server.
```

```
==== End of Episode Summary ====
```

```
Steps Taken: 30/30
```

```
Discounted Episode Reward: -7.185280246516709
```

```
=====
```

Corresponding Server-side logs:

Episode 1:

```
Moved to state ID: 18 with Reward: -0.02
DEBUG (Server): stopMethod=0
User agent1 reached maximum steps per episode.
Attempting to retrieve user: agent1 with memory reference: 1225264877
Action Taken: RIGHT, New State: 18, Reward: -0.02
Current users in userMap:
- agent1 (equals provided username? true)
Sent GAME_ACTION_REWARD_RESPONSE with action: 3, reward: -0.02, nextStateId: 18
User found: agent1
Sent GAME_FINAL_STATE_RESPONSE
```

```
End of Episode Summary:
- Total Episodes: 1
- Successful Episodes: 0
- Steps Taken: 30
- Discounted Episode Reward: -8.811421888987708
```

## Episode 2:

```
Moved to state ID: 22 with Reward: -0.02
DEBUG (Server): stopMethod=0
User agent1 reached maximum steps per episode.
Action Taken: DOWN, New State: 22, Reward: -0.02
Sent GAME_ACTION_REWARD_RESPONSE with action: 1, reward: -0.02, nextStateId: 22
Sent GAME_FINAL_STATE_RESPONSE
```

```
End of Episode Summary:
- Total Episodes: 2
- Successful Episodes: 0
- Steps Taken: 30
- Discounted Episode Reward: -11.213227822475048
```

## Episode 3:

```
Moved to state ID: 16 with Reward: -0.02
DEBUG (Server): stopMethod=0
User agent1 reached maximum steps per episode.
Action Taken: RIGHT, New State: 16, Reward: -0.02
Sent GAME_ACTION_REWARD_RESPONSE with action: 3, reward: -0.02, nextStateId: 16
Sent GAME_FINAL_STATE_RESPONSE
```

```
End of Episode Summary:
- Total Episodes: 3
- Successful Episodes: 0
- Steps Taken: 30
- Discounted Episode Reward: -7.185280246516709
```

STOP\_METHOD=1:

Client-side logs:

## Episode 1:

```
Received EXTENSION_RESPONSE: cmd=rl.action
Received messageType: GAME_ACTION_REWARD_RESPONSE
Parsing messageType: GAME_ACTION_REWARD_RESPONSE
All received keys: [reward, messageType, nextStateId, action]
```



```
State transition: 19 -> 24
!!! GOAL STATE REACHED !!!
```

```
=== Episode 1 Information ===
```

```
Reason: Goal state reached!
```

```
Goal Reached: Yes!
```

```
=====
```

```
=== End of Episode Summary ===
```

```
Steps Taken: 52/30
```

```
Discounted Episode Reward: -6.861007529129636
```

Episode 2:

```
Received EXTENSION_RESPONSE: cmd=rl.action
```

```
Received messageType: GAME_ACTION_REWARD_RESPONSE
```

```
Parsing messageType: GAME_ACTION_REWARD_RESPONSE
```

```
All received keys: [reward, messageType, nextStateId, action]
```

```
State transition: 23 -> 24
```

```
!!! GOAL STATE REACHED !!!
```

```
=== Episode 2 Information ===
```

```
Reason: Goal state reached!
```

```
Goal Reached: Yes!
```

```
=====
```

```
=== End of Episode Summary ===
```

```
Steps Taken: 120/30
```

```
Discounted Episode Reward: -12.303766160414735
```

Episode 3:

```
Received EXTENSION_RESPONSE: cmd=rl.action
```

```
Received messageType: GAME_ACTION_REWARD_RESPONSE
```

```
Parsing messageType: GAME_ACTION_REWARD_RESPONSE
```

```
All received keys: [reward, messageType, nextStateId, action]
```

```
State transition: 23 -> 24
```

```
!!! GOAL STATE REACHED !!!
```

```
=== Episode 3 Information ===
```

```
Reason: Goal state reached!
```

```
Goal Reached: Yes!
```

```
=====
```

```
=== End of Episode Summary ===
```

```
Steps Taken: 104/30
```

```
Discounted Episode Reward: -0.39777134658552254
```

Corresponding Server-side logs:

Episode 1:

```
Moved to state ID: 24 with Reward: 20.0
```

```
DEBUG (Server): stopMethod=1
```

```
User agent1 has reached the terminal state: 24
```

```
Action Taken: DOWN, New State: 24, Reward: 20.0
```

```
Sent GAME_ACTION_REWARD_RESPONSE with action: 1, reward: 20.0, nextStateId: 24
Sent GAME_FINAL_STATE_RESPONSE
```

End of Episode Summary:

- Total Episodes: 1
- Successful Episodes: 0
- Steps Taken: 52
- Discounted Episode Reward: -6.861007529129636

Episode 2:

```
Moved to state ID: 24 with Reward: 20.0
```

```
DEBUG (Server): stopMethod=1
```

```
User agent1 has reached the terminal state: 24
```

```
Action Taken: RIGHT, New State: 24, Reward: 20.0
```

```
Sent GAME_ACTION_REWARD_RESPONSE with action: 3, reward: 20.0, nextStateId: 24
```

```
Sent GAME_FINAL_STATE_RESPONSE
```

End of Episode Summary:

- Total Episodes: 2
- Successful Episodes: 0
- Steps Taken: 120
- Discounted Episode Reward: -12.303766160414735

Episode 3:

```
Moved to state ID: 24 with Reward: 20.0
```

```
DEBUG (Server): stopMethod=1
```

```
User agent1 has reached the terminal state: 24
```

```
Action Taken: RIGHT, New State: 24, Reward: 20.0
```

```
Sent GAME_ACTION_REWARD_RESPONSE with action: 3, reward: 20.0, nextStateId: 24
```

```
Sent GAME_FINAL_STATE_RESPONSE
```

End of Episode Summary:

- Total Episodes: 3
- Successful Episodes: 0
- Steps Taken: 104
- Discounted Episode Reward: -0.39777134658552254

STOP\_METHOD=2:

Client-side logs:

Episode 1:

```
State transition: 6 -> 7
```

```
DEBUG: stopMethod=2
```

```
Updated Q-value for state 6, action 3: 0.5995598535615114
```

```
Sent Q-Table updates to the server.
```

```
Sent V-Table updates to the server.
```

```
Epsilon decayed to: 0.946354579813443
```

```
Received EXTENSION_RESPONSE: cmd=r1.action
```

```
Received messageType: GAME_FINAL_STATE_RESPONSE
```

```
Parsing messageType: GAME_FINAL_STATE_RESPONSE
```

```
All received keys: [stepsThisEpisode, messageType, successfulEpisodes, cumulativeReward, totalEpisodes, isTerminal]
```

```
=== End of Episode Summary ===
```

```
Steps Taken: 11/30
```

```
Discounted Episode Reward: -1.189491060782
```

## Episode 2:

```
State transition: 15 -> 20
DEBUG: stopMethod=2
Updated Q-value for state 15, action 1: 0.6100182199622346
Sent Q-Table updates to the server.
Sent V-Table updates to the server.
Epsilon decayed to: 0.8911090557802088
Received EXTENSION_RESPONSE: cmd=rl.action
Received messageType: GAME_FINAL_STATE_RESPONSE
Parsing messageType: GAME_FINAL_STATE_RESPONSE
All received keys: [stepsThisEpisode, messageType, successfulEpisodes, cumulativeReward, totalEpisodes, isTerminal]

=== End of Episode Summary ===
Steps Taken: 12/30
Discounted Episode Reward: -3.5293140927038005
```

## Episode 3:

```
State transition: 2 -> 1
DEBUG: stopMethod=2
Updated Q-value for state 2, action 2: 0.5287384634894934
Sent Q-Table updates to the server.
Sent V-Table updates to the server.
Epsilon decayed to: 0.851801859600347
Received EXTENSION_RESPONSE: cmd=rl.action
Received messageType: GAME_FINAL_STATE_RESPONSE
Parsing messageType: GAME_FINAL_STATE_RESPONSE
All received keys: [stepsThisEpisode, messageType, successfulEpisodes, cumulativeReward, totalEpisodes, isTerminal]

=== End of Episode Summary ===
Steps Taken: 9/30
Discounted Episode Reward: -6.4488139022
```

## Corresponding Server-side logs:

### Episode 1:

```
Moved to state ID: 7 with Reward: -0.02
DEBUG (Server): stopMethod=2
User agent1 stopped due to STOP_METHOD=2 random stopping condition.
Action Taken: RIGHT, New State: 7, Reward: -0.02
Sent GAME_ACTION_REWARD_RESPONSE with action: 3, reward: -0.02, nextStateId: 7
Sent GAME_FINAL_STATE_RESPONSE

End of Episode Summary:
- Total Episodes: 1
- Successful Episodes: 0
- Steps Taken: 11
- Discounted Episode Reward: -1.189491060782
```

### Episode 2:

```
Moved to state ID: 20 with Reward: -0.02
User found: agent1
DEBUG (Server): stopMethod=2
```

```
User agent1 stopped due to STOP_METHOD=2 random stopping condition.  
Action Taken: DOWN, New State: 20, Reward: -0.02  
Sent GAME_ACTION_REWARD_RESPONSE with action: 1, reward: -0.02, nextStateId: 20  
Sent GAME_FINAL_STATE_RESPONSE
```

End of Episode Summary:

- Total Episodes: 2
- Successful Episodes: 0
- Steps Taken: 12
- Discounted Episode Reward: -3.5293140927038005

Episode 3:

```
Moved to state ID: 1 with Reward: -0.02  
DEBUG (Server): stopMethod=2  
User agent1 stopped due to STOP_METHOD=2 random stopping condition.  
Action Taken: LEFT, New State: 1, Reward: -0.02  
Sent GAME_ACTION_REWARD_RESPONSE with action: 2, reward: -0.02, nextStateId: 1  
Sent GAME_FINAL_STATE_RESPONSE
```

End of Episode Summary:

- Total Episodes: 3
- Successful Episodes: 0
- Steps Taken: 9
- Discounted Episode Reward: -6.4488139022

As one can see from the logs for all the stop methods, they stop exactly as described above in the stopping method descriptions.

### References

1. Sutton, Richard S., and Andrew G. Barto. Reinforcement learning: An introduction. MIT press, 2018.
2. <https://www.smartfoxserver.com/download/sfs2x#p=installer> – SmartFoxServer Java Server API downloads page
3. <https://www.smartfoxserver.com/download/sfs2x#p=client> - SmartFoxServer Java Client API downloads page
4. <https://www.oracle.com/ca-en/java/technologies/downloads/> - Oracle Java Downloads page
5. <https://www.geeksforgeeks.org/epsilon-greedy-algorithm-in-reinforcement-learning/> - Epsilon-Greedy policy in RL